

Parallel Computations on Immutable Data Structures

Florian Biermann
fbie@itu.dk

IT University of Copenhagen & UCAS

2016-06-02



IT University
of Copenhagen


A Recursive Function

$$\sum_{i=1}^6 i = (+\ 1\ (+\ 2\ (+\ 3\ (+\ 4\ (+\ 5\ 6))))))$$

A Recursive Function

$$\sum_{i=1}^6 i = (+\ 1\ (+\ 2\ (+\ 3\ (+\ 4\ (+\ 5\ 6))))))$$

The expression's call-tree:



What Is The Problem?

All operations are performed sequentially, even though it would be perfectly ok to compute

`(+ 1 (+ 2 3))`

and

`(+ 4 (+ 5 6))`

in parallel.

What Is The Problem?

All operations are performed sequentially, even though it would be perfectly ok to compute

`(+ 1 (+ 2 3))`

and

`(+ 4 (+ 5 6))`

in parallel.

We can re-formulate the expression and obtain the same result!


A Parallel Expression

$$\sum_{i=1}^6 i = (+ (+ 1 (+ 2 3)) (+ 4 (+ 5 6)))$$

A Parallel Expression

$$\sum_{i=1}^6 i = (+ (+ 1 (+ 2 3)) (+ 4 (+ 5 6)))$$


The expression's call-tree:



A Parallel Expression

$$\sum_{i=1}^6 i = (+ (+ 1 (+ 2 3)) (+ 4 (+ 5 6)))$$

The expression's call-tree:




Why is this possible? Because + is an associative operation!

$$(+ \text{ } x \text{ } (+ \text{ } y \text{ } z)) = (+ \text{ } (+ \text{ } x \text{ } y) \text{ } z)$$

Generalizing Parallel Expressions

We can do this for any function f iff

$$(f\ a\ (f\ b\ c)) = (f\ (f\ a\ b)\ c)$$



Another Problem

Mapping a function `f` over a `cons` list.


```
(: seq-map (All (A B)
  (-> (-> A B) (Listof A) (Listof B))))
(define (seq-map f xs)
  (match xs
    ['() '()]
    [(cons x xs) (cons (f x) (seq-map f xs))]))
```

Another Problem

Mapping a function `f` over a `cons` list.

```
(: seq-map (All (A B)
  (-> (-> A B) (Listof A) (Listof B))))
(define (seq-map f xs)
  (match xs
    ['() '()]
    [(cons x xs) (cons (f x) (seq-map f xs))]))
```

Call-tree of `seq-map`:



What Is The Problem?

`seq-map` is not the problem. It is the only way to handle a `cons` list. **The real problem is the single-linked list!** It is inherently sequential.

What Is The Problem?

`seq-map` is not the problem. It is the only way to handle a `cons` list. **The real problem is the single-linked list!** It is inherently sequential.

All parallel expressions we saw so far were trees. Maybe we can develop a tree data structure that we can use to implement a parallel version of `map`?

A New Parallel Data Type

We represent the list as a tree instead:

```
(define-type (CatListof A) (U (leaf A) (cat A)))  
(struct (A) leaf ([a : A]))  
(struct (A) cat ([l : (CatListof A)]  
                 [r : (CatListof A)]))
```

A New Parallel Data Type

We represent the list as a tree instead:

```
(define-type (CatListof A) (U (leaf A) (cat A)))  
(struct (A) leaf ([a : A]))  
(struct (A) cat ([l : (CatListof A)]  
                 [r : (CatListof A)]))
```

(leaf a) produces a singleton `CatList`.

(cat l r) concatenates two `CatList` instances.

Let's Implement Map!

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))
(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r) (cat (par-map f l)
                     (par-map f r))]))
```


Let's Implement Map!

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))
(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r) (cat (par-map f l)
                     (par-map f r))]))
```

Is this efficient? There are a few things to be aware of:

- ▶ A list has $O(n)$ memory overhead, a tree $O(n \log n)$.
- ▶ If the tree is not balanced, we lose parallelism.
- ▶ It is not truly parallel yet. This is the next step!

Parallelism in Racket

In Racket we use a fork/join style model of parallel computations. Fork is called `future` and join is called `touch`:

```
(: in-parallel (All (A B C)
  (-> (-> A B) (-> A C) A (Pairof B C))))
(define (in-parallel f g x)
  (let ([fx (future (lambda () (f x)))]
        [gx (g x)]))
    (cons (touch fx) gx)))
```

Parallelism in Racket

In Racket we use a fork/join style model of parallel computations. Fork is called `future` and join is called `touch`:

```
(: in-parallel (All (A B C)
  (-> (-> A B) (-> A C) A (Pairof B C))))
(define (in-parallel f g x)
  (let ([fx (future (lambda () (f x)))]
        [gx (g x)]))
    (cons (touch fx) gx)))
```

Racket has a built-in scheduler for `futures`, so we give a lot of control to the run-time.

```
(: future (All (A) (->(->A) (Futureof A))))
(: touch (All (A) (->(Futureof A) A)))
```

Parallelizing Map, For Real

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))
(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r)
     (let ([l0 (future ;; Map l in parallel.
                     (lambda () (par-map f l)))]
           [r0 (future ;; Map r in parallel.
                     (lambda () (par-map f r)))]])
       (cat (touch l0) (touch r0)))]))
```