

Parallel Computations on Immutable Data Structures

Florian Biermann
fbie@itu.dk

IT University of Copenhagen & UCAS

2016-06-02



IT University
of Copenhagen


A Sequential Expression

(+ 1 (+ 2 (+ 3 (+ 4 (+ 5 6)))))

A Sequential Expression

`(+ 1 (+ 2 (+ 3 (+ 4 (+ 5 6)))))`

The expression tree:



What Is The Problem?

All operations are performed sequentially, even though it would be perfectly ok to compute

`(+ 1 (+ 2 3))`

and

`(+ 4 (+ 5 6))`

in parallel. The result would be the same.

What Is The Problem?

All operations are performed sequentially, even though it would be perfectly ok to compute

`(+ 1 (+ 2 3))`

and

`(+ 4 (+ 5 6))`

in parallel. The result would be the same.

We can re-formulate the expression and obtain the same result!

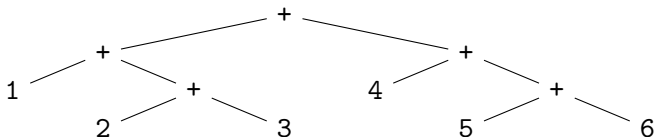
A Parallel Expression

`(+ (+ 1 (+ 2 3)) (+ 4 (+ 5 6)))`

A Parallel Expression

$(+ (+ 1 (+ 2 3)) (+ 4 (+ 5 6)))$

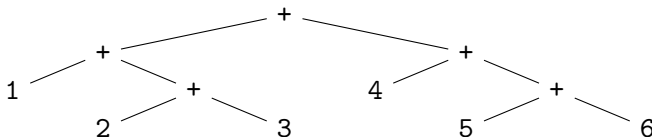
The expression tree:



A Parallel Expression

$(+ (+ 1 (+ 2 3)) (+ 4 (+ 5 6)))$

The expression tree:



Why is this possible? Because + is an associative operation!

$$(+ \ x \ (+ \ y \ z)) = (+ \ (+ \ x \ y) \ z)$$

Generalizing Parallel Expressions


We can do this for any function f iff

$$(f\ a\ (f\ b\ c)) = (f\ (f\ a\ b)\ c)$$

Generalizing Parallel Expressions

We can do this for any function f iff

$$(f\ a\ (f\ b\ c)) = (f\ (f\ a\ b)\ c)$$



Another Problem

Mapping a function `f` over a `cons` list.

```
(: seq-map (All (A B)
  (-> (-> A B) (Listof A) (Listof B))))
```

Another Problem

Mapping a function `f` over a `cons` list.

```
(: seq-map (All (A B)
  (-> (-> A B) (Listof A) (Listof B))))

(define (seq-map f xs)
  (match xs
    ['() '()]
    [(cons x xs) (cons (f x) (seq-map f xs))]))
```


Another Problem

Mapping a function `f` over a `cons` list.

```
(: seq-map (All (A B)
  (-> (-> A B) (Listof A) (Listof B))))

(define (seq-map f xs)
  (match xs
    ['() '()]
    [(cons x xs) (cons (f x) (seq-map f xs))]))
```

Expression tree of `seq-map`:

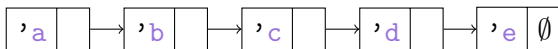


What Is The Problem?

`seq-map` is not the problem. It is the only way to handle a `cons` list. **The real problem is the single-linked list!** It is inherently sequential:

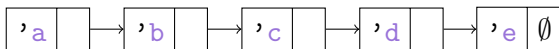
What Is The Problem?

`seq-map` is not the problem. It is the only way to handle a `cons` list. **The real problem is the single-linked list!** It is inherently sequential:



What Is The Problem?

`seq-map` is not the problem. It is the only way to handle a `cons` list. **The real problem is the single-linked list!** It is inherently sequential:



The parallel expressions we have seen so far are *trees*. Maybe we can develop a tree data structure that we can use to implement a parallel version of `map`?

A New Parallel Data Type

We represent the list as a tree instead:

```
(define-type (CatListof A) (U (leaf A)
                               (cat A)))

(struct (A) leaf ([a : A]))
(struct (A) cat  ([l : (CatListof A)]
                  [r : (CatListof A)]))
```

A New Parallel Data Type

We represent the list as a tree instead:

```
(define-type (CatListof A) (U (leaf A)
                               (cat A)))

(struct (A) leaf ([a : A]))
(struct (A) cat  ([l : (CatListof A)]
                  [r : (CatListof A)]))
```

(leaf a) produces a singleton CatList.

(cat l r) concatenates two CatList instances.

A New Parallel Data Type

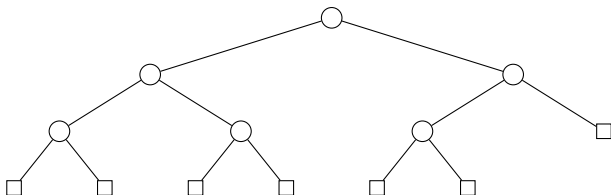
We represent the list as a tree instead:

```
(define-type (CatListof A) (U (leaf A)
                               (cat A)))

(struct (A) leaf ([a : A]))
(struct (A) cat  ([l : (CatListof A)]
                  [r : (CatListof A)]))
```

(leaf a) produces a singleton CatList.

(cat l r) concatenates two CatList instances.



Let's Implement Map!

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))
```

Let's Implement Map!

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))

(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r) (cat (par-map f l)
                     (par-map f r))]))
```

Let's Implement Map!

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))

(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r) (cat (par-map f l)
                     (par-map f r))]))
```

Note: This function is not *truly* parallel yet. We have just created a *possibility* for parallelism.

Parallelism in Racket

In Racket we use a fork/join style model of parallel computations. Fork is called `future` and join is called `touch`:

```
(: in-parallel (All (A B C)
  (-> (-> A B) (-> A C) A (Pairof B C))))
```

Parallelism in Racket

In Racket we use a fork/join style model of parallel computations. Fork is called `future` and join is called `touch`:

```
(: in-parallel (All (A B C)
  (-> (-> A B) (-> A C) A (Pairof B C))))
```

```
(define (in-parallel f g x)
  (let ([fx (future (lambda () (f x)))]
        [gx (g x)]))
    (cons (touch fx) gx)))
```


Parallelism in Racket

In Racket we use a fork/join style model of parallel computations. Fork is called `future` and join is called `touch`:

```
(: in-parallel (All (A B C)
  (-> (-> A B) (-> A C) A (Pairof B C))))
```

```
(define (in-parallel f g x)
  (let ([fx (future (lambda () (f x)))]
        [gx (g x)]))
    (cons (touch fx) gx)))
```

Racket has a built-in scheduler for `futures`, so we give a lot of control to the run-time.

```
(: future (All (A) (->(->A) (Futureof A))))
(: touch (All (A) (->(Futureof A) A)))
```

Parallelizing Map, For Real

```
(: par-map (All (A B)  
  (-> (-> A B) (CatListof A) (CatListof B))))
```

Parallelizing Map, For Real

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))

(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))])
```

Parallelizing Map, For Real

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))

(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r)
     (let ([fl (future ;; Map l in parallel.
                      (lambda () (par-map f l)))]
```

Parallelizing Map, For Real

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))

(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r)
     (let ([fl (future ;; Map l in parallel.
                     (lambda () (par-map f l)))]
           [fr (future ;; Map r in parallel.
                     (lambda () (par-map f r)))]))])])
```

Parallelizing Map, For Real

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))

(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r)
     (let ([fl (future ;; Map l in parallel.
                      (lambda () (par-map f l)))]
           [fr (future ;; Map r in parallel.
                      (lambda () (par-map f r)))]])
       (cat (touch fl) (touch fr)))]))
```

The recursive call is wrapped in a `future`!

Pro and Contra

The good parts:

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.
- ▶ A thread that is blocked by `touch` is de-scheduled until the `touched future` is done, no locking, spinning etc.

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.
- ▶ A thread that is blocked by `touch` is de-scheduled until the `touched future` is done, no locking, spinning etc.
- ▶ As a result, we do not have to manually decide **how** to split up work across hardware threads.

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.
- ▶ A thread that is blocked by `touch` is de-scheduled until the `touched future` is done, no locking, spinning etc.
- ▶ As a result, we do not have to manually decide **how** to split up work across hardware threads.

The bad parts:

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.
- ▶ A thread that is blocked by `touch` is de-scheduled until the `touched future` is done, no locking, spinning etc.
- ▶ As a result, we do not have to manually decide **how** to split up work across hardware threads.

The bad parts:

- ▶ A list has $O(n)$ memory overhead, a balanced tree $O(n \log n)$.

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.
- ▶ A thread that is blocked by `touch` is de-scheduled until the `touched future` is done, no locking, spinning etc.
- ▶ As a result, we do not have to manually decide **how** to split up work across hardware threads.

The bad parts:

- ▶ A list has $O(n)$ memory overhead, a balanced tree $O(n \log n)$.
- ▶ If the tree is not balanced, we lose parallelism – mitigate by balancing algorithm!

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.
- ▶ A thread that is blocked by `touch` is de-scheduled until the `touched future` is done, no locking, spinning etc.
- ▶ As a result, we do not have to manually decide **how** to split up work across hardware threads.

The bad parts:

- ▶ A list has $O(n)$ memory overhead, a balanced tree $O(n \log n)$.
- ▶ If the tree is not balanced, we lose parallelism – mitigate by balancing algorithm!
- ▶ Maybe we spawn **excessively many futures** and put too much strain on the scheduler?

An Improved Variant

If the size of the tree is too small, then the overhead of constructing futures dominates the run-time cost. We can handle this problem by using lists as leaves instead of scalar values.

An Improved Variant

If the size of the tree is too small, then the overhead of constructing `future`s dominates the run-time cost. We can handle this problem by using lists as leaves instead of scalar values.

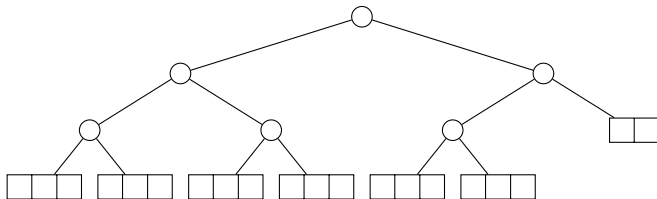
We call such a tree a **Rope**:

```
(define-type (Ropeof A) (U (leaf A)
                           (cat A)))

(struct (A) leaf ([as : (Listof A)]))
(struct (A) cat  ([l  : (Ropeof A)]
                  [r  : (Ropeof A)]))
```


Ropes: Trees With Lists as Leaves

We must define a **maximum size** s_{\max} for leaves. In this example, $s_{\max} = 3$:



Parallel higher-order functions on ropes.

Live Coding Session

Parallel higher-order functions on ropes.

- ▶ I code on the big screen and show you how to do functional programming in Racket.

Live Coding Session

Parallel higher-order functions on ropes.

- ▶ I code on the big screen and show you how to do functional programming in Racket.
- ▶ You help with ideas and suggestions and experiment yourselves!

Live Coding Session

Parallel higher-order functions on ropes.

- ▶ I code on the big screen and show you how to do functional programming in Racket.
- ▶ You help with ideas and suggestions and experiment yourselves!
- ▶ All code we write will be available on
 - ▶ github.com/fbie/parallel-functional-programming

Live Coding Session

Parallel higher-order functions on ropes.

- ▶ I code on the big screen and show you how to do functional programming in Racket.
- ▶ You help with ideas and suggestions and experiment yourselves!
- ▶ All code we write will be available on
 - ▶ github.com/fbie/parallel-functional-programming
- ▶ Want to start now and code along?
 - ▶ Download the code from [2/ropes-lecture.rkt](#)

Live Coding Session

Parallel higher-order functions on ropes.

- ▶ I code on the big screen and show you how to do functional programming in Racket.
- ▶ You help with ideas and suggestions and experiment yourselves!
- ▶ All code we write will be available on
 - ▶ github.com/fbie/parallel-functional-programming
- ▶ Want to start now and code along?
 - ▶ Download the code from [2/ropes-lecture.rkt](#)
- ▶ You can download Racket from
 - ▶ racket-lang.org

New Syntax: Recursive Let

Standard `let`:

```
(let ([x 1]
      [y (+ 2 x)]))
  (* y x))
```


New Syntax: Recursive Let

Standard `let`:

```
(let ([x 1]
      [y (+ 2 x)]))
  (* y x))
```

Undefined identifier: `x`

New Syntax: Recursive Let

Standard `let`:

```
(let ([x 1]
      [y (+ 2 x)]))
  (* y x))
```

Undefined identifier: `x`

Solution: `letrec`, recursive let-binding:

```
(letrec ([x 1]
         [y (+ 2 x)]))
  (* y x))
```

Allows bindings to reference one-another.

De-sugaring Recursive Let

```
(letrec ([x 1]
         [y (+ 2 x)])
  (* y x))
```

is equal to

```
(let ([x 1])
  (let ([y (+ 2 x)])
    (* y x)))
```

Take-Aways From This Lecture

Take-Aways From This Lecture

- ▶ Immutable data structures are automatically thread-safe!

Take-Aways From This Lecture

- ▶ Immutable data structures are automatically thread-safe!
- ▶ Choose trees over lists.

Take-Aways From This Lecture

- ▶ Immutable data structures are automatically thread-safe!
- ▶ Choose trees over lists.
- ▶ Use sequential data structures at the leaves.

Take-Aways From This Lecture

- ▶ Immutable data structures are automatically thread-safe!
- ▶ Choose trees over lists.
- ▶ Use sequential data structures at the leaves.
- ▶ Use the built-in scheduler.

Take-Aways From This Lecture

- ▶ Immutable data structures are automatically thread-safe!
- ▶ Choose trees over lists.
- ▶ Use sequential data structures at the leaves.
- ▶ Use the built-in scheduler.
- ▶ Use higher-order functions for re-usable solutions.

Take-Aways From This Lecture

- ▶ Immutable data structures are automatically thread-safe!
- ▶ Choose trees over lists.
- ▶ Use sequential data structures at the leaves.
- ▶ Use the built-in scheduler.
- ▶ Use higher-order functions for re-usable solutions.

Reading material available on the GitHub page:

Take-Aways From This Lecture

- ▶ Immutable data structures are automatically thread-safe!
- ▶ Choose trees over lists.
- ▶ Use sequential data structures at the leaves.
- ▶ Use the built-in scheduler.
- ▶ Use higher-order functions for re-usable solutions.

Reading material available on the GitHub page:

- ▶ Herb Sutter: **The Free Lunch is Over, Dr. Dobb's Journal**

Take-Aways From This Lecture

- ▶ Immutable data structures are automatically thread-safe!
- ▶ Choose trees over lists.
- ▶ Use sequential data structures at the leaves.
- ▶ Use the built-in scheduler.
- ▶ Use higher-order functions for re-usable solutions.

Reading material available on the GitHub page:

- ▶ Herb Sutter: **The Free Lunch is Over, Dr. Dobb's Journal**
- ▶ Michael Erdmann: **Parallelism, Cost Graphs, and Sequences, lecture notes**

Take-Aways From This Lecture

- ▶ Immutable data structures are automatically thread-safe!
- ▶ Choose trees over lists.
- ▶ Use sequential data structures at the leaves.
- ▶ Use the built-in scheduler.
- ▶ Use higher-order functions for re-usable solutions.

Reading material available on the GitHub page:

- ▶ Herb Sutter: **The Free Lunch is Over, Dr. Dobb's Journal**
- ▶ Michael Erdmann: **Parallelism, Cost Graphs, and Sequences, lecture notes**
- ▶ Hans-J. Boehm et al.: **Ropes: an Alternative to Strings** (optional)

Questions?

Slides and code available at
github.com/fbie/parallel-functional-programming.

Questions?

Slides and code available at
github.com/fbie/parallel-functional-programming.

Any questions? E-mail me at fbie@itu.dk.

Questions?

Slides and code available at
github.com/fbie/parallel-functional-programming.

Any questions? E-mail me at fbie@itu.dk.

Thank you for your attention!