

Parallel Computations on Immutable Data Structures

Florian Biermann
fbie@itu.dk

IT University of Copenhagen & UCAS

2016-06-02



IT University
of Copenhagen


A Sequential Expression

$$\sum_{i=1}^6 i = (+\ 1\ (+\ 2\ (+\ 3\ (+\ 4\ (+\ 5\ 6))))))$$

A Sequential Expression

$$\sum_{i=1}^6 i = (+\ 1\ (+\ 2\ (+\ 3\ (+\ 4\ (+\ 5\ 6))))))$$

The expression tree:



What Is The Problem?

All operations are performed sequentially, even though it would be perfectly ok to compute

`(+ 1 (+ 2 3))`

and

`(+ 4 (+ 5 6))`

in parallel. The result would be the same.

What Is The Problem?

All operations are performed sequentially, even though it would be perfectly ok to compute

`(+ 1 (+ 2 3))`

and

`(+ 4 (+ 5 6))`

in parallel. The result would be the same.

We can re-formulate the expression and obtain the same result!

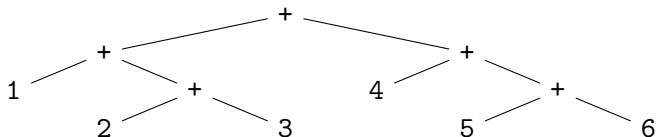
A Parallel Expression

$$\sum_{i=1}^6 i = (+ (+ 1 (+ 2 3)) (+ 4 (+ 5 6)))$$

A Parallel Expression

$$\sum_{i=1}^6 i = (+ (+ 1 (+ 2 3)) (+ 4 (+ 5 6)))$$

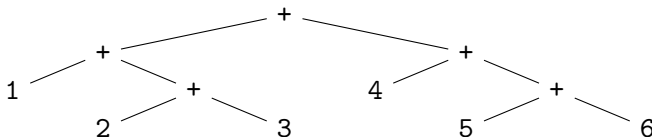
The expression tree:



A Parallel Expression

$$\sum_{i=1}^6 i = (+ (+ 1 (+ 2 3)) (+ 4 (+ 5 6)))$$

The expression tree:



Why is this possible? Because + is an associative operation!

$$(+ \text{ x } (+ \text{ y } \text{ z})) = (+ (+ \text{ x } \text{ y}) \text{ z})$$

Generalizing Parallel Expressions


We can do this for any function f iff

$$(f\ a\ (f\ b\ c)) = (f\ (f\ a\ b)\ c)$$

Generalizing Parallel Expressions

We can do this for any function f iff

$$(f\ a\ (f\ b\ c)) = (f\ (f\ a\ b)\ c)$$



Another Problem

Mapping a function `f` over a `cons` list.

```
(: seq-map (All (A B)
  (-> (-> A B) (Listof A) (Listof B))))
```

Another Problem

Mapping a function `f` over a `cons` list.

```
(: seq-map (All (A B)
  (-> (-> A B) (Listof A) (Listof B))))

(define (seq-map f xs)
  (match xs
    ['() '()]
    [(cons x xs) (cons (f x) (seq-map f xs))]))
```


Another Problem

Mapping a function `f` over a `cons` list.

```
(: seq-map (All (A B)
  (-> (-> A B) (Listof A) (Listof B))))

(define (seq-map f xs)
  (match xs
    ['() '()]
    [(cons x xs) (cons (f x) (seq-map f xs))]))
```

Expression tree of `seq-map`:

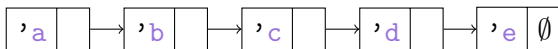


What Is The Problem?

`seq-map` is not the problem. It is the only way to handle a `cons` list. **The real problem is the single-linked list!** It is inherently sequential:


What Is The Problem?

`seq-map` is not the problem. It is the only way to handle a `cons` list. **The real problem is the single-linked list!** It is inherently sequential:



What Is The Problem?

`seq-map` is not the problem. It is the only way to handle a `cons` list. **The real problem is the single-linked list!** It is inherently sequential:



The parallel expressions we have seen so far are *trees*. Maybe we can develop a tree data structure that we can use to implement a parallel version of `map`?

A New Parallel Data Type

We represent the list as a tree instead:

```
(define-type (CatListof A) (U (leaf A)
                               (cat A)))

(struct (A) leaf ([a : A]))
(struct (A) cat  ([l : (CatListof A)]
                  [r : (CatListof A)]))
```

A New Parallel Data Type

We represent the list as a tree instead:

```
(define-type (CatListof A) (U (leaf A)
                               (cat A)))

(struct (A) leaf ([a : A]))
(struct (A) cat  ([l : (CatListof A)]
                  [r : (CatListof A)]))
```

(leaf a) produces a singleton CatList.

(cat l r) concatenates two CatList instances.

A New Parallel Data Type

We represent the list as a tree instead:


```
(define-type (CatListof A) (U (leaf A)
                               (cat A)))

(struct (A) leaf ([a : A]))

(struct (A) cat ([l : (CatListof A)]
                 [r : (CatListof A)]))
```

(leaf a) produces a singleton CatList.

(cat l r) concatenates two CatList instances.



Let's Implement Map!

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))
```

Let's Implement Map!

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))

(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r) (cat (par-map f l)
                     (par-map f r))]))
```

Let's Implement Map!

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))

(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r) (cat (par-map f l)
                     (par-map f r))]))
```

Note: This function is not *truly* parallel yet. We have just created a *possibility* for parallelism.

Parallelism in Racket

In Racket we use a fork/join style model of parallel computations. Fork is called `future` and join is called `touch`:

```
(: in-parallel (All (A B C)
  (-> (-> A B) (-> A C) A (Pairof B C))))
```

Parallelism in Racket

In Racket we use a fork/join style model of parallel computations. Fork is called `future` and join is called `touch`:

```
(: in-parallel (All (A B C)
  (-> (-> A B) (-> A C) A (Pairof B C))))
```

```
(define (in-parallel f g x)
  (let ([fx (future (lambda () (f x)))]
        [gx (g x)]))
    (cons (touch fx) gx)))
```


Parallelism in Racket

In Racket we use a fork/join style model of parallel computations. Fork is called `future` and join is called `touch`:

```
(: in-parallel (All (A B C)
  (-> (-> A B) (-> A C) A (Pairof B C))))
```

```
(define (in-parallel f g x)
  (let ([fx (future (lambda () (f x)))]
        [gx (g x)]))
    (cons (touch fx) gx)))
```

Racket has a built-in scheduler for `futures`, so we give a lot of control to the run-time.

```
(: future (All (A) (->(->A) (Futureof A))))
(: touch (All (A) (->(Futureof A) A)))
```

Parallelizing Map, For Real

```
(: par-map (All (A B)  
  (-> (-> A B) (CatListof A) (CatListof B))))
```

Parallelizing Map, For Real

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))

(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))])
```

Parallelizing Map, For Real

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))

(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r)
     (let ([fl (future ;; Map l in parallel.
                      (lambda () (par-map f l)))]
```

Parallelizing Map, For Real

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))

(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r)
     (let ([fl (future ;; Map l in parallel.
                      (lambda () (par-map f l)))]
           [fr (future ;; Map r in parallel.
                      (lambda () (par-map f r)))]])
```

Parallelizing Map, For Real

```
(: par-map (All (A B)
  (-> (-> A B) (CatListof A) (CatListof B))))

(define (par-map f xs)
  (match xs
    [(leaf x) (leaf (f x))]
    [(cat l r)
     (let ([fl (future ;; Map l in parallel.
                      (lambda () (par-map f l)))]
           [fr (future ;; Map r in parallel.
                      (lambda () (par-map f r)))]])
       (cat (touch fl) (touch fr)))]))
```

The recursive call is wrapped in a `future`!

Pro and Contra

The good parts:

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.
- ▶ A thread that is blocked by `touch` is de-scheduled until the `touched future` is done, no locking, spinning etc.

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.
- ▶ A thread that is blocked by `touch` is de-scheduled until the `touched future` is done, no locking, spinning etc.
- ▶ As a result, we do not have to manually decide **how** to split up work across hardware threads.

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.
- ▶ A thread that is blocked by `touch` is de-scheduled until the `touched future` is done, no locking, spinning etc.
- ▶ As a result, we do not have to manually decide **how** to split up work across hardware threads.

The bad parts:

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.
- ▶ A thread that is blocked by `touch` is de-scheduled until the `touched future` is done, no locking, spinning etc.
- ▶ As a result, we do not have to manually decide **how** to split up work across hardware threads.

The bad parts:

- ▶ A list has $O(n)$ memory overhead, a balanced tree $O(n \log n)$.

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.
- ▶ A thread that is blocked by `touch` is de-scheduled until the `touched future` is done, no locking, spinning etc.
- ▶ As a result, we do not have to manually decide **how** to split up work across hardware threads.

The bad parts:

- ▶ A list has $O(n)$ memory overhead, a balanced tree $O(n \log n)$.
- ▶ If the tree is not balanced, we lose parallelism – mitigate by balancing algorithm!

Pro and Contra

The good parts:

- ▶ The build-in scheduler decides on the number of threads to use.
- ▶ A thread that is blocked by `touch` is de-scheduled until the `touched future` is done, no locking, spinning etc.
- ▶ As a result, we do not have to manually decide **how** to split up work across hardware threads.

The bad parts:

- ▶ A list has $O(n)$ memory overhead, a balanced tree $O(n \log n)$.
- ▶ If the tree is not balanced, we lose parallelism – mitigate by balancing algorithm!
- ▶ Maybe we spawn **excessively many futures** and put too much strain on the scheduler?

An Improved Variant

If the size of the tree is too small, then the overhead of constructing `future`s dominates the run-time cost. We can handle this problem by using lists as leaves instead of scalar values.


We call such a tree a **Rope**:

```
(define-type (Ropeof A) (U (leaf A)
                           (cat A)))

(struct (A) leaf ([as : (Listof A)]))
(struct (A) cat  ([l  : (Ropeof A)]
                  [r  : (Ropeof A)]))
```

Ropes: Trees With Lists as Leaves

We must define a **maximum size** s_{\max} for leaves. In this example, $s_{\max} = 3$:



Parallel higher-order functions on ropes.

Live Coding Session

Parallel higher-order functions on ropes.

- ▶ I code on the big screen and show you how to do functional programming in Racket.

Live Coding Session

Parallel higher-order functions on ropes.

- ▶ I code on the big screen and show you how to do functional programming in Racket.
- ▶ You help with ideas and suggestions and experiment yourselves!

Live Coding Session

Parallel higher-order functions on ropes.

- ▶ I code on the big screen and show you how to do functional programming in Racket.
- ▶ You help with ideas and suggestions and experiment yourselves!
- ▶ All code we write will be available on
 - ▶ github.com/fbie/parallel-functional-programming

Live Coding Session

Parallel higher-order functions on ropes.

- ▶ I code on the big screen and show you how to do functional programming in Racket.
- ▶ You help with ideas and suggestions and experiment yourselves!
- ▶ All code we write will be available on
 - ▶ github.com/fbie/parallel-functional-programming
- ▶ Want to start now and code along?
 - ▶ Download the code from [2/ropes-lecture.rkt](#)

Live Coding Session

Parallel higher-order functions on ropes.

- ▶ I code on the big screen and show you how to do functional programming in Racket.
- ▶ You help with ideas and suggestions and experiment yourselves!
- ▶ All code we write will be available on
 - ▶ github.com/fbie/parallel-functional-programming
- ▶ Want to start now and code along?
 - ▶ Download the code from [2/ropes-lecture.rkt](#)
- ▶ You can download Racket from
 - ▶ racket-lang.org

New Syntax: Recursive Let

Standard `let`:

```
(let ([x 1]
      [y (+ 2 x)]))
  (* y x))
```

New Syntax: Recursive Let

Standard `let`:

```
(let ([x 1]
      [y (+ 2 x)]))
  (* y x))
```

Undefined identifier: `x`

New Syntax: Recursive Let

Standard `let`:

```
(let ([x 1]
      [y (+ 2 x)]))
  (* y x))
```

Undefined identifier: `x`

Solution: `letrec`, recursive let-binding:

```
(letrec ([x 1]
         [y (+ 2 x)]))
  (* y x))
```

Allows bindings to reference one-another.

De-sugaring Recursive Let

```
(letrec ([x 1]
         [y (+ 2 x)])
  (* y x))
```

is equal to

```
(let ([x 1])
  (let ([y (+ 2 x)])
    (* y x)))
```